

```
In [1]: import pandas as pd
import numpy as np

import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px

from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.metrics import confusion_matrix, classification_report
```

```
In [2]: df = pd.read_csv('heart_v2.csv')
```

```
In [3]: df.head()
```

Out[3]:

	age	sex	BP	cholesterol	heart disease
0	70	1	130	322	1
1	67	0	115	564	0
2	57	1	124	261	1
3	64	1	128	263	0
4	74	0	120	269	0

```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 270 entries, 0 to 269
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0   age             270 non-null   int64
1   sex             270 non-null   int64
2   BP              270 non-null   int64
3   cholesterol     270 non-null   int64
4   heart disease   270 non-null   int64
dtypes: int64(5)
memory usage: 10.7 KB
```

```
In [5]: df.columns
```

Out[5]: Index(['age', 'sex', 'BP', 'cholesterol', 'heart disease'], dtype='object')

```
In [6]: df.describe()
```

```
Out[6]:
```

	age	sex	BP	cholesterol	heart disease
count	270.000000	270.000000	270.000000	270.000000	270.000000
mean	54.433333	0.677778	131.344444	249.659259	0.444444
std	9.109067	0.468195	17.861608	51.686237	0.497827
min	29.000000	0.000000	94.000000	126.000000	0.000000
25%	48.000000	0.000000	120.000000	213.000000	0.000000
50%	55.000000	1.000000	130.000000	245.000000	0.000000
75%	61.000000	1.000000	140.000000	280.000000	1.000000
max	77.000000	1.000000	200.000000	564.000000	1.000000

```
In [7]: df.isnull().sum()
```

```
Out[7]: age          0
sex          0
BP          0
cholesterol  0
heart disease 0
dtype: int64
```


```
In [8]: df['heart disease'].value_counts()
```

```
Out[8]: 0    150
1     120
Name: heart disease, dtype: int64
```


```
In [9]: df.shape
```

```
Out[9]: (270, 5)
```

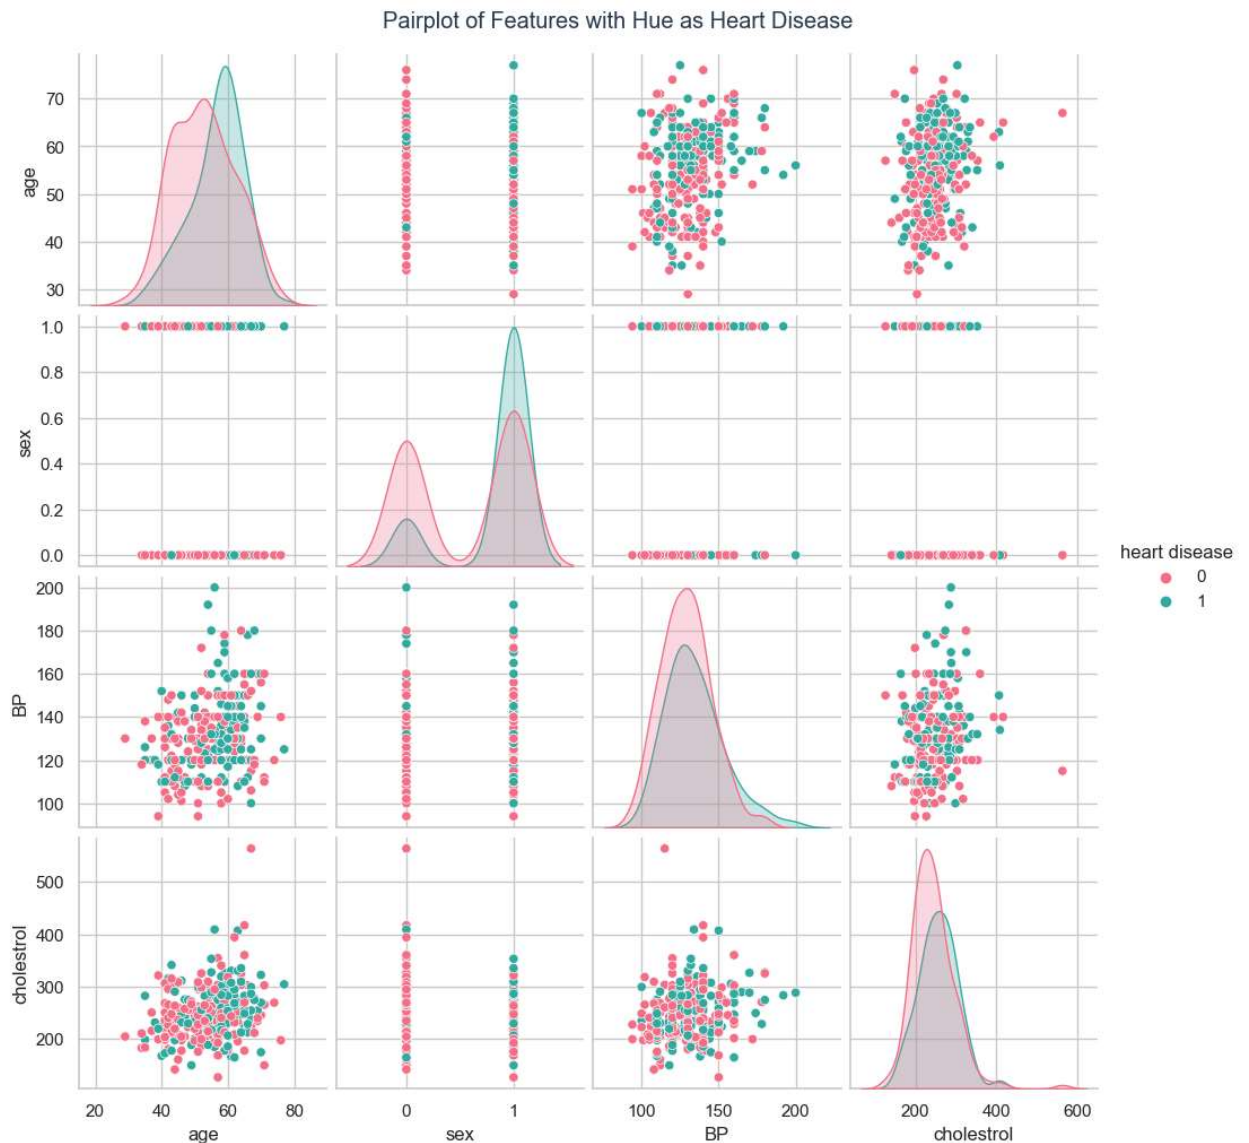
```
In [10]: fig = px.histogram(df, x='age', nbins=20, title='Distribution of Age', labels={'age': 'Age', 'count': 'count'})  
fig.show()
```



```
In [11]: # Blood Pressure Distribution
fig = px.bar(df, x='BP', title='Distribution of Blood Pressure', color='BP', template='
fig.show()
```



```
In [12]: sns.set_theme(style="whitegrid")
sns.pairplot(df, hue='heart disease', palette='husl')
plt.suptitle('Pairplot of Features with Hue as Heart Disease', y=1.02, color='#2c3e50')
plt.show()
```



```
In [13]: # Putting feature variable to X
X = df.drop('heart disease',axis=1)

# Putting response variable to y
y = df['heart disease']
```

```
In [14]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, train_size=0.7, random_state=
```

```
In [15]: X_train.shape, X_test.shape
```

```
Out[15]: ((189, 4), (81, 4))
```

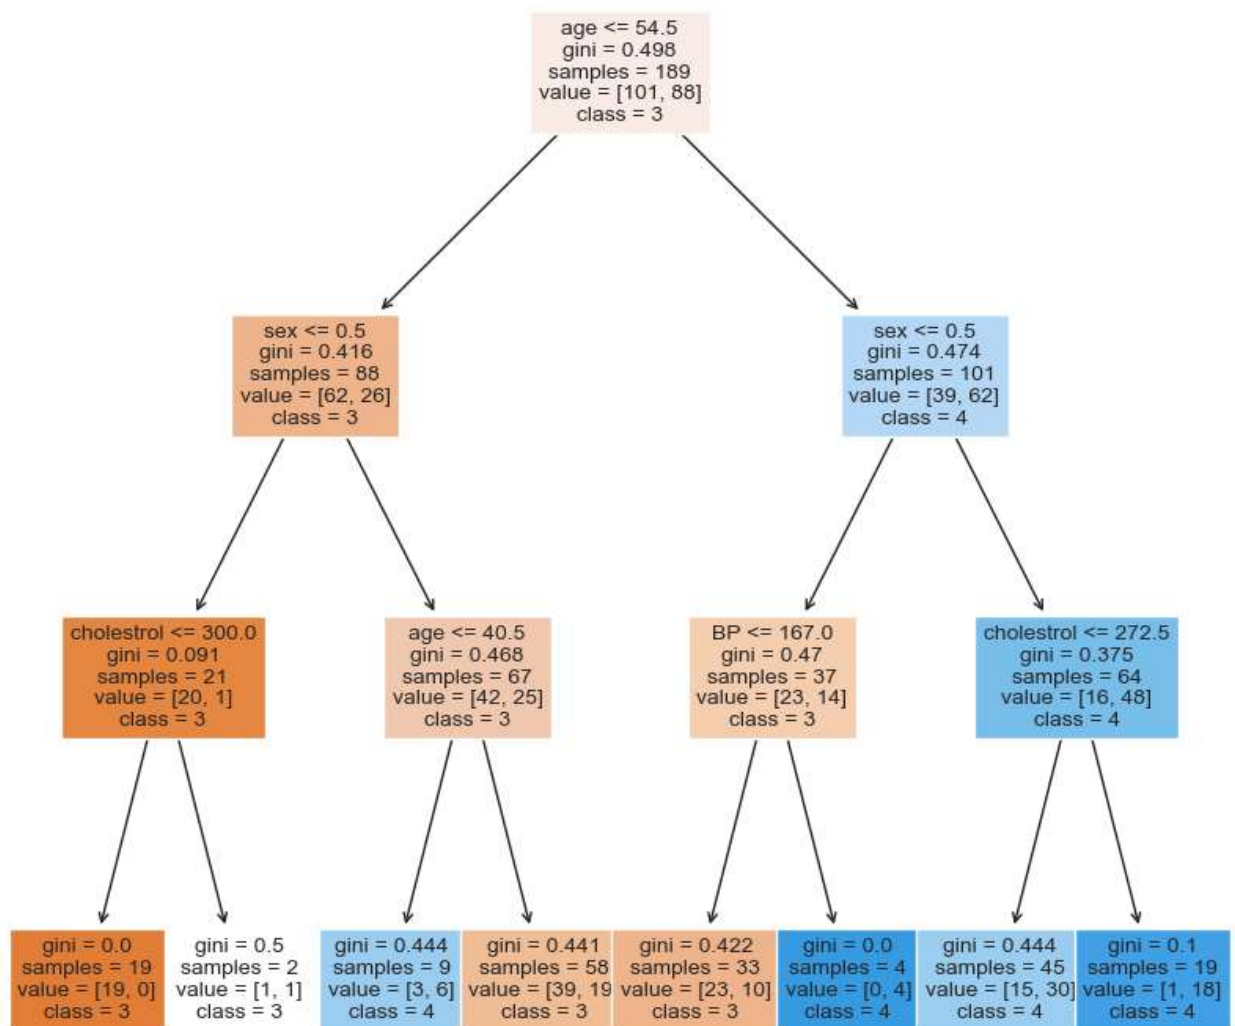
```
In [16]: from sklearn.tree import DecisionTreeClassifier, plot_tree
dt = DecisionTreeClassifier(max_depth=3)
dt.fit(X_train, y_train)
```

```
Out[16]: DecisionTreeClassifier
DecisionTreeClassifier(max_depth=3)
```

(<https://scikit-learn.org/1.4/modules/generated/sklearn.tree.DecisionTreeClassifier>)

```
In [17]: plt.figure(figsize=(10, 10))
plot_tree(dt, filled=True, feature_names=X.columns, class_names=['3', '4', '5', '6', '7'])
plt.title("Decision Tree Visualization for Red Wine Quality", fontsize=16)
plt.show()
```

## Decision Tree Visualization for Red Wine Quality



```
In [18]: y_train = dt.predict(X_train)
y_pred = dt.predict(X_test)
```

```
In [19]: from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
class_report = classification_report(y_test, y_pred)
print(class_report)
```

	precision	recall	f1-score	support
0	0.66	0.71	0.69	49
1	0.50	0.44	0.47	32
accuracy			0.60	81
macro avg	0.58	0.58	0.58	81
weighted avg	0.60	0.60	0.60	81

```
In [20]: print(accuracy_score(y_test, y_pred))
```


```
0.6049382716049383
```

```
In [21]: matrix=confusion_matrix(y_test,y_pred)
print(matrix)
```

```
[[35 14]
 [18 14]]
```

```
In [22]: plt.figure(figsize = (10,7))
sns.heatmap(matrix, annot=True, cmap='inferno', linewidths=.5, fmt='g')
```

Out[22]: <Axes: >



## Mushroom dataset

```
In [23]: #Load data
df = pd.read_csv("mushrooms.csv")
```



In [24]:

```
df.head()
```

Out[24]:

	class	cap- shape	cap- surface	cap- color	bruises	odor	gill- attachment	gill- spacing	gill- size	gill- color	...	stalk- surface- below- ring	stalk- color- above- ring	stal colo below rir
0	p	x	s	n	t	p	f	c	n	k	...	s	w	
1	e	x	s	y	t	a	f	c	b	k	...	s	w	
2	e	b	s	w	t	l	f	c	b	n	...	s	w	
3	p	x	y	w	t	p	f	c	n	n	...	s	w	
4	e	x	s	g	f	n	f	w	b	k	...	s	w	

5 rows × 23 columns



In [25]:

```
df.isnull().sum()
```

Out[25]:

```
class                0
cap-shape            0
cap-surface          0
cap-color            0
bruises              0
odor                 0
gill-attachment      0
gill-spacing         0
gill-size            0
gill-color           0
stalk-shape          0
stalk-root           0
stalk-surface-above-ring 0
stalk-surface-below-ring 0
stalk-color-above-ring 0
stalk-color-below-ring 0
veil-type            0
veil-color           0
ring-number          0
ring-type            0
spore-print-color    0
population           0
habitat              0
dtype: int64
```

In [26]:

```
df['class'].unique()
```

Out[26]:


```
array(['p', 'e'], dtype=object)
```

```
In [27]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8124 entries, 0 to 8123
Data columns (total 23 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   class                                8124 non-null   object
1   cap-shape                            8124 non-null   object
2   cap-surface                          8124 non-null   object
3   cap-color                           8124 non-null   object
4   bruises                             8124 non-null   object
5   odor                                 8124 non-null   object
6   gill-attachment                      8124 non-null   object
7   gill-spacing                        8124 non-null   object
8   gill-size                           8124 non-null   object
9   gill-color                          8124 non-null   object
10  stalk-shape                         8124 non-null   object
11  stalk-root                          8124 non-null   object
12  stalk-surface-above-ring            8124 non-null   object
13  stalk-surface-below-ring            8124 non-null   object
14  stalk-color-above-ring              8124 non-null   object
15  stalk-color-below-ring              8124 non-null   object
16  veil-type                           8124 non-null   object
17  veil-color                          8124 non-null   object
18  ring-number                         8124 non-null   object
19  ring-type                           8124 non-null   object
20  spore-print-color                   8124 non-null   object
21  population                          8124 non-null   object
22  habitat                             8124 non-null   object
dtypes: object(23)
memory usage: 1.4+ MB
```

```
In [28]: sns.histplot(df['class'])
```

```
Out[28]: <Axes: xlabel='class', ylabel='Count'>
```



## Seprating Features and Targets:

```
In [29]: X = df.drop('class',axis=1)
y = df['class']
```

```
In [30]: X = pd.get_dummies(X)
X.head()
```

```
Out[30]:
```

	cap- shape_b	cap- shape_c	cap- shape_f	cap- shape_k	cap- shape_s	cap- shape_x	cap- surface_f	cap- surface_g	cap- surface_s	cap- surface_y	...
0	0	0	0	0	0	1	0	0	1	0	...
1	0	0	0	0	0	1	0	0	1	0	...
2	1	0	0	0	0	0	0	0	1	0	...
3	0	0	0	0	0	1	0	0	0	1	...
4	0	0	0	0	0	1	0	0	1	0	...

5 rows × 117 columns



## Label Encoding

```
In [31]: from sklearn.preprocessing import LabelEncoder
encoder = LabelEncoder()
y = encoder.fit_transform(y)
print(y)
```

```
[1 0 0 ... 0 1 0]
```

## Splitting into training and testing:

```
In [32]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=1)
```

## Creating Decision Tree using entropy:

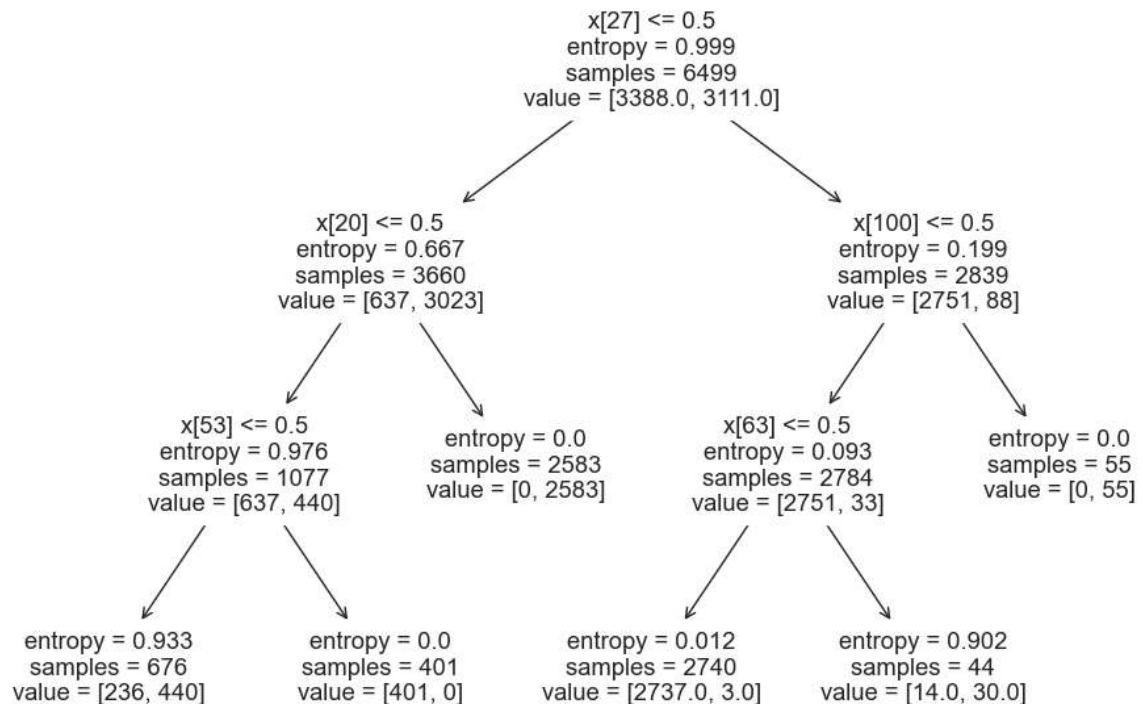
```
In [33]: clf = DecisionTreeClassifier(criterion='entropy', max_depth=3, random_state=0)
clf.fit(X_train, y_train)
```

```
Out[33]: DecisionTreeClassifier
DecisionTreeClassifier(criterion='entropy', max_depth=3, random_state=0)
```

<https://scikit-learn.org/1.4/module>

```
In [34]: from sklearn import tree
plt.figure(figsize=(12,8))
tree.plot_tree(clf.fit(X_train, y_train))
```

```
Out[34]: [Text(0.5555555555555556, 0.875, 'x[27] <= 0.5\nentropy = 0.999\nsamples = 6499\nvalue = [3388.0, 3111.0]'),
Text(0.3333333333333333, 0.625, 'x[20] <= 0.5\nentropy = 0.667\nsamples = 3660\nvalue = [637, 3023]'),
Text(0.2222222222222222, 0.375, 'x[53] <= 0.5\nentropy = 0.976\nsamples = 1077\nvalue = [637, 440]'),
Text(0.1111111111111111, 0.125, 'entropy = 0.933\nsamples = 676\nvalue = [236, 440]'),
Text(0.3333333333333333, 0.125, 'entropy = 0.0\nsamples = 401\nvalue = [401, 0]'),
Text(0.4444444444444444, 0.375, 'entropy = 0.0\nsamples = 2583\nvalue = [0, 2583]'),
Text(0.7777777777777778, 0.625, 'x[100] <= 0.5\nentropy = 0.199\nsamples = 2839\nvalue = [2751, 88]'),
Text(0.6666666666666666, 0.375, 'x[63] <= 0.5\nentropy = 0.093\nsamples = 2784\nvalue = [2751, 33]'),
Text(0.5555555555555556, 0.125, 'entropy = 0.012\nsamples = 2740\nvalue = [2737.0, 3.0]'),
Text(0.7777777777777778, 0.125, 'entropy = 0.902\nsamples = 44\nvalue = [14.0, 30.0]'),
Text(0.8888888888888888, 0.375, 'entropy = 0.0\nsamples = 55\nvalue = [0, 55]')]
```



```
In [35]: #Predict values
y_pred = clf.predict(X_test)
```

```
In [36]: #Predict values using x_train
y_pred_train = clf.predict(X_train)
```

## Calculating accuracy\_score from scikit\_learn

```
In [37]: print('criterion entropy accuracy: {0:.2f}'.format(accuracy_score(y_test, y_pred)*100))
print('Training set: {0:.2f}'.format(accuracy_score(y_train, y_pred_train)*100))
```

```
criterion entropy accuracy: 96.37
Training set: 96.11
```

## Calculating accuracy\_score from model of the classifier

```
In [38]: print('Training set score: {0:.2f}'.format(clf.score(X_train, y_train)*100))
print('Test set score: {0:.2f}'.format(clf.score(X_test, y_test)*100))
```


```
Training set score: 96.11
Test set score: 96.37
```

```
In [39]: cm = confusion_matrix(y_test, y_pred)
print(cm)
```

```
[[766  54]
 [  5 800]]
```

In [40]:

```
plt.subplots(figsize=(10, 10))
sns.heatmap(cm, annot=True, linewidths=0.5,fmt= '.0f')
plt.show()
```



In [41]: `print(classification_report(y_test, y_pred))`

	precision	recall	f1-score	support
0	0.99	0.93	0.96	820
1	0.94	0.99	0.96	805
accuracy			0.96	1625
macro avg	0.97	0.96	0.96	1625
weighted avg	0.97	0.96	0.96	1625

```
In [42]: from sklearn.metrics import f1_score  
f1_score = f1_score(y_test, y_pred)  
print(f1_score)
```

0.9644364074743822

In [ ]: